

HW1: Classification

Jiafeng Chen

Yufeng Ling

Francisco Rivera

February 4, 2019

1 Introduction

This note lays out our expectation for a homework submission in *CS287: Statistical Natural Language Processing*. While you do not have to follow this template to the letter, we do expect that write-ups have a very clear structure and cover all the elements described in this note. With this in mind, the burden is on the presenter to demonstrate why deviations from the standard are necessary.

All write-ups should include a short introduction. In this section you should summarize the underlying problem in high-level language and describe the extensions that you have decided to propose in your implementation. When you describe these extensions you should carefully cite the papers of interest. For instance, it will often be useful to cite the work seen in class ([Murphy, 2012](#)). Alternatively, you can also cite papers inline, for instance the work of [Berger et al. \(1996\)](#).

2 Problem Description—Old

In general, homeworks will be specified using informal language. As part of the assignment, we expect you to write-out a definition of the problem and your model in formal language. For this class, we will use the following notation:

- \mathbf{b}, \mathbf{m} ; bold letters for vectors.
- \mathbf{B}, \mathbf{M} ; bold capital letters for matrices.
- \mathcal{B}, \mathcal{M} ; script-case for sets.
- b_i, x_i ; lower case for scalars or indexing into vectors.

For instance in natural language processing, it is common to use discrete sets like \mathcal{V} for the vocabulary of the language, or \mathcal{T} for a tag set of the language. We might also want

one-hot vectors representing words. These will be of the type $v \in \{0, 1\}^{|\mathcal{V}|}$. In a note, it is crucial to define the types of all variables that are introduced. The problem description is the right place to do this.

3 Problem Description

The focus of the problem set is sentiment classification. We are given *sentences* and aim to classify them as either positive or negative sentiment (a binary classification). These predictions can be made in a probabilistic fashion by writing y_i as the event that the i th sentence is positive sentiment, and calculating $p(y_i)$.

Sentences are themselves sequences of words $w \in \mathcal{V}$ in some vocabulary \mathcal{V} . In particular, there will be two special words, $w_{\text{unk}}, w_{\text{pad}} \in \mathcal{V}$ which represent an unknown word and a padding unit respectively; we address their significance later. A particular sequence of words w_1, \dots, w_n need not have a specific length n , which varies across sentences.

We represent words in two main ways. The first is as a one-hot encoded vector $v \in \{0, 1\}^{|\mathcal{V}|}$. This representation is useful for the **Logistic Regression** model. Alternatively, we can use a dense embedding. That is, each word gets assigned a vector $v \in \mathbb{R}^d$ where d is the embedding dimension. We use [Mikolov et al. \(2013\)](#)'s pre-trained word embeddings ($d = 300$) for the **Continuous Bag of Words** and **Convolutional Neural Network** models.

4 Model and Algorithms

4.1 Naive Bayes

4.2 Logistic Regression

4.3 Continuous Bag of Words

A downside of logistic regression is that we have as many parameters as the size of our vocabulary \mathcal{V} . This means that without an extensive training set, we run the risk of overfitting. To address this, we can reduce the dimensionality of our embedding by using [Mikolov et al. \(2013\)](#)'s word embeddings. This model has three parts:

1. We start by using average-pooling over time on the word-embeddings to get a vector of dimension d for each sentence.
2. Then, we pass that vector through a linear transform to get another vector of dimension d_2 which then gets passed through a non-linear transformation (ReLU) to make up the hidden layer.

3. Finally, we pass the hidden layer output through another linear layer to get a vector with two elements, and a softmax transformation to turn these values into the probabilities of positive and negative sentiment.

The model has two parameter matrices of size $d \times d_2$ and $d_2 \times 2$ respectively. We implement the model for $d_2 = 100$. We train using Adam for 20 epochs. In summary, the model's steps are visualized in [Figure 1](#).

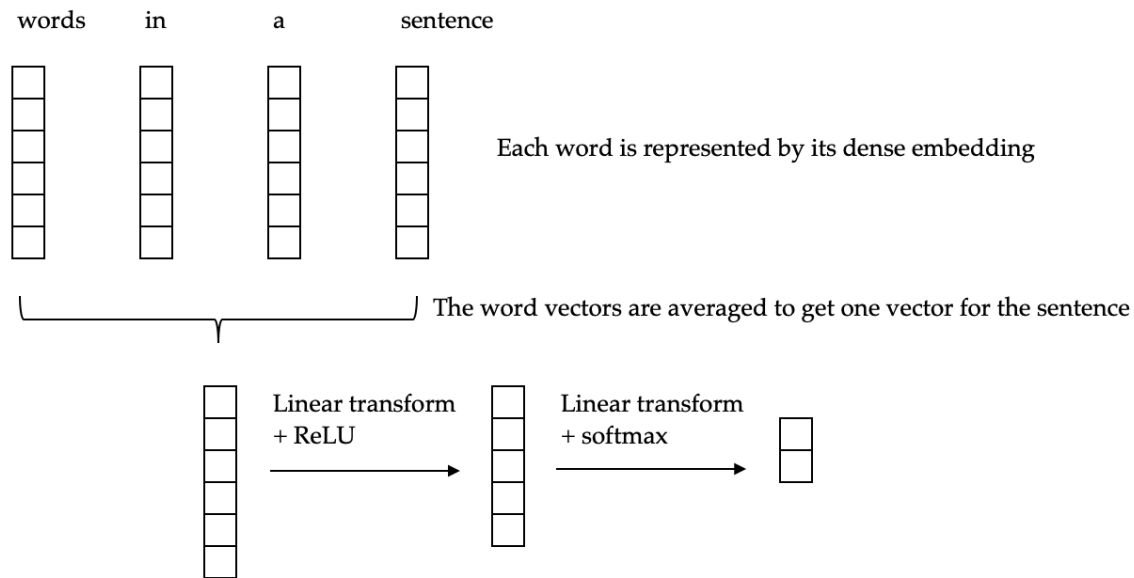


Figure 1: CBOW model pipeline

We call this model a *bag of words* because the order of the words is lost in the average-pooling over time. In other words, any permutation of a sentence will result in the same prediction by this model.

4.4 Convolutional Neural Network

While the [Continuous Bag of Words](#) model makes use of the pre-trained dense embeddings, it suffers a limitation from the bag-of-words construction. Because the order of words is not taken into account, the sentences “the movie was good, not boring” and “the movie was boring, not good” cannot be distinguished.

5 Experiments

Finally we end with the experimental section. Each assignment will make clear the main experiments and baselines that you should run. For these experiments you should present

a main results table. Here we give a sample Table 1. In addition to these results you should describe in words what the table shows and the relative performance of the models.

Besides the main results we will also ask you to present other results comparing particular aspects of the models. For instance, for word embedding experiments, we may ask you to show a chart of the projected word vectors. This experiment will lead to something like Figure 3. This should also be described within the body of the text itself.

Model	Acc.
BASELINE 1	0.45
BASELINE 2	2.59
MODEL 1	10.59
MODEL 2	13.42
MODEL 3	7.49

Table 1: Table with the main results.

6 Conclusion

End the write-up with a very short recap of the main experiments and the main results. Describe any challenges you may have faced, and what could have been improved in the model.

References

- Berger, A. L., Pietra, V. J. D., and Pietra, S. A. D. (1996). A maximum entropy approach to natural language processing. *Computational linguistics*, 22(1):39–71.
- Mikolov, T., Chen, K., Corrado, G., and Dean, J. (2013). Efficient estimation of word representations in vector space. *arXiv preprint arXiv:1301.3781*.
- Murphy, K. P. (2012). *Machine learning: a probabilistic perspective*. MIT press.

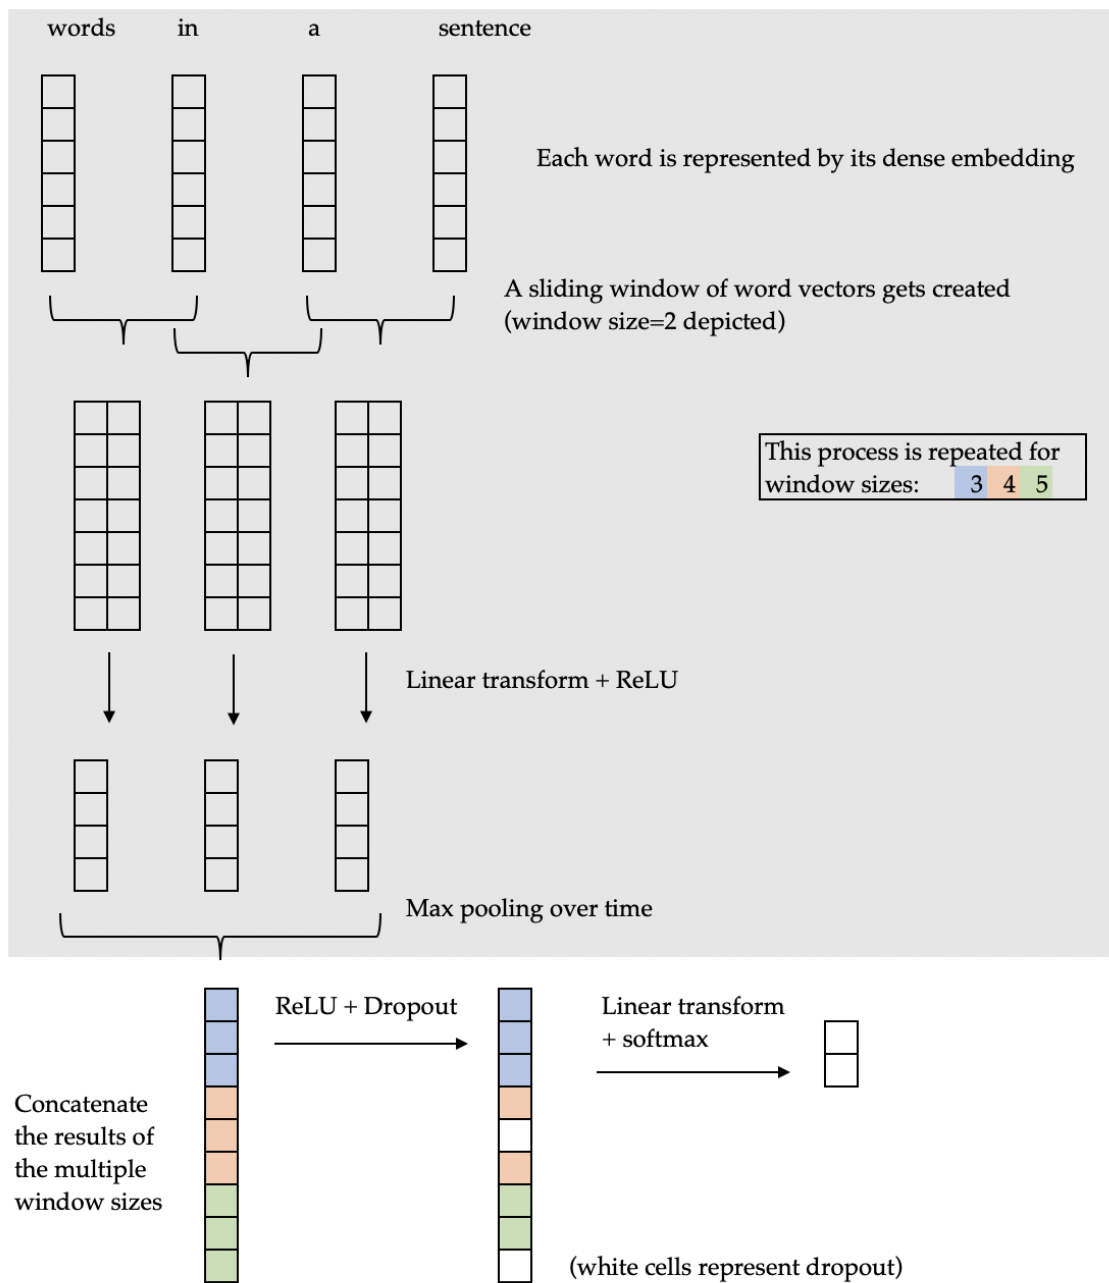


Figure 2: Convolutional Neural Network Model



Figure 3: Sample qualitative chart.